

Submit Apache Spark Applications Lab



Objectives

In this lab, you will:

- Install a Spark Master and Worker using Docker Compose
- Create a python script containing a spark job
- Submit the job to the cluster directly from python (Note: you'll learn how to submit a job from the command line in the Kubernetes Lab)

Prerequisites

Note: If you are running this lab within the Skillsnetwork Lab environment, all prerequisites are already installed for you

The only pre-requisites to this lab are:

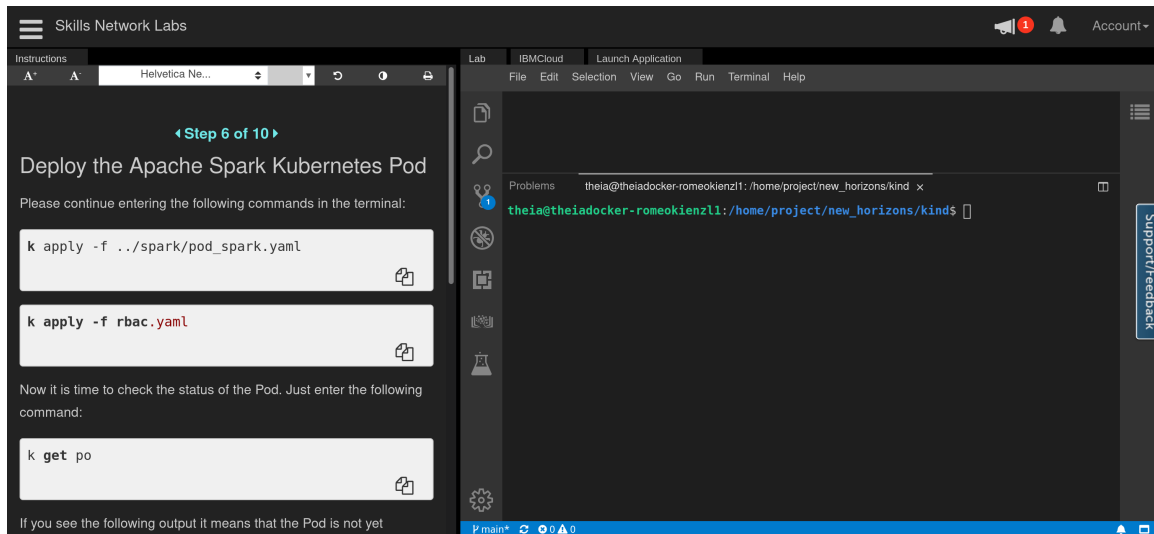
- A working *docker* installation
- Docker Compose
- The *git* command line tool
- A python development environment

Project Overview

Welcome to the lab on how to submit Apache Spark applications from a python script. This exercise is straightforward thanks to Docker Compose.

Install a Apache Spark cluster using Docker Compose

On the right hand side to this instructions you'll see the Theia IDE. Select the *Lab* tab. On the menu bar select *Terminal*>*New Terminal*.



Please enter the following commands in the terminal:
Install *PySpark*:

1. 1
1. `python3 -m pip install pyspark`

Copied!

Get the latest code:

1. 1
1. `git clone https://github.com/big-data-europe/docker-spark.git`

Copied!

Change the directory to the downloaded code:

1. 1
1. `cd docker-spark`

Copied!

Start the cluster

1. 1
1. `docker-compose up`


Copied!

After quite some time you should see the following message:
Successfully registered with master spark://<server address>:7077

```
spark-worker-1 | 22/01/14 05:39:36 INFO Worker: Successfully registered with  
master spark://50bfae057469:7077
```

Create code

Create a new python file called *submit.py*



Paste the following code to the file and save.

```
1. 1  
2. 2  
3. 3  
4. 4  
5. 5  
6. 6  
7. 7  
8. 8  
9. 9  
10. 10  
11. 11  
12. 12  
13. 13  
14. 14  
15. 15  
16. 16  
17. 17  
18. 18  
19. 19  
20. 20  
21. 21
```

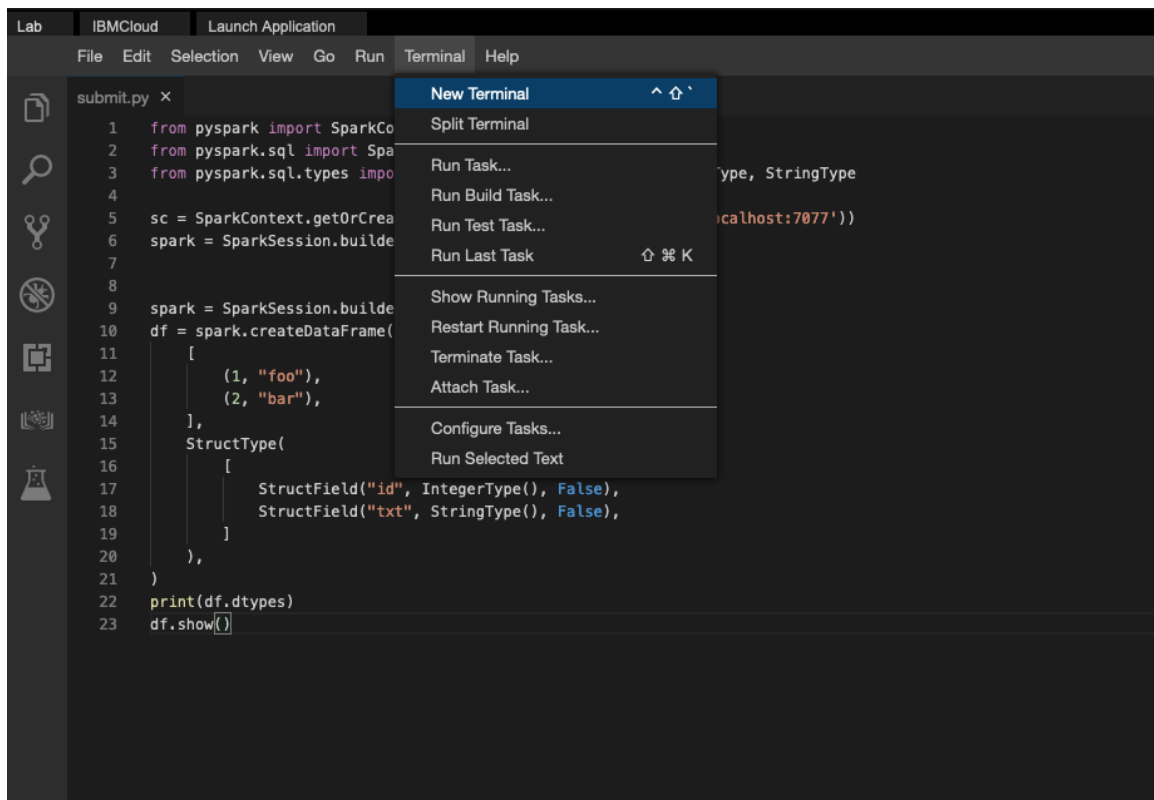
```
22. 22
23. 23
24. 24
```

```
1. from pyspark import SparkContext, SparkConf
2. from pyspark.sql import SparkSession
3. from pyspark.sql.types import StructField, StructType, IntegerType, StringType
4.
5. sc = SparkContext.getOrCreate(SparkConf().setMaster('spark://localhost:7077'))
6. sc.setLogLevel("INFO")
7.
8. spark = SparkSession.builder.getOrCreate()
9.
10. spark = SparkSession.builder.getOrCreate()
11. df = spark.createDataFrame(
12.     [
13.         (1, "foo"),
14.         (2, "bar"),
15.     ],
16.     StructType(
17.         [
18.             StructField("id", IntegerType(), False),
19.             StructField("txt", StringType(), False),
20.         ]
21.     ),
22. )
23. print(df.dtypes)
24. df.show()
```

Copied!

Execute code / submit Spark job

Now we execute the python file we saved earlier. Click on “Terminal” on the bar on top and click on “New Terminal” as shown in the image below.



Once the terminal opens up at the bottom of the window, please type in the following command in the terminal to execute the Python script:

1. 1

1. python3 submit.py

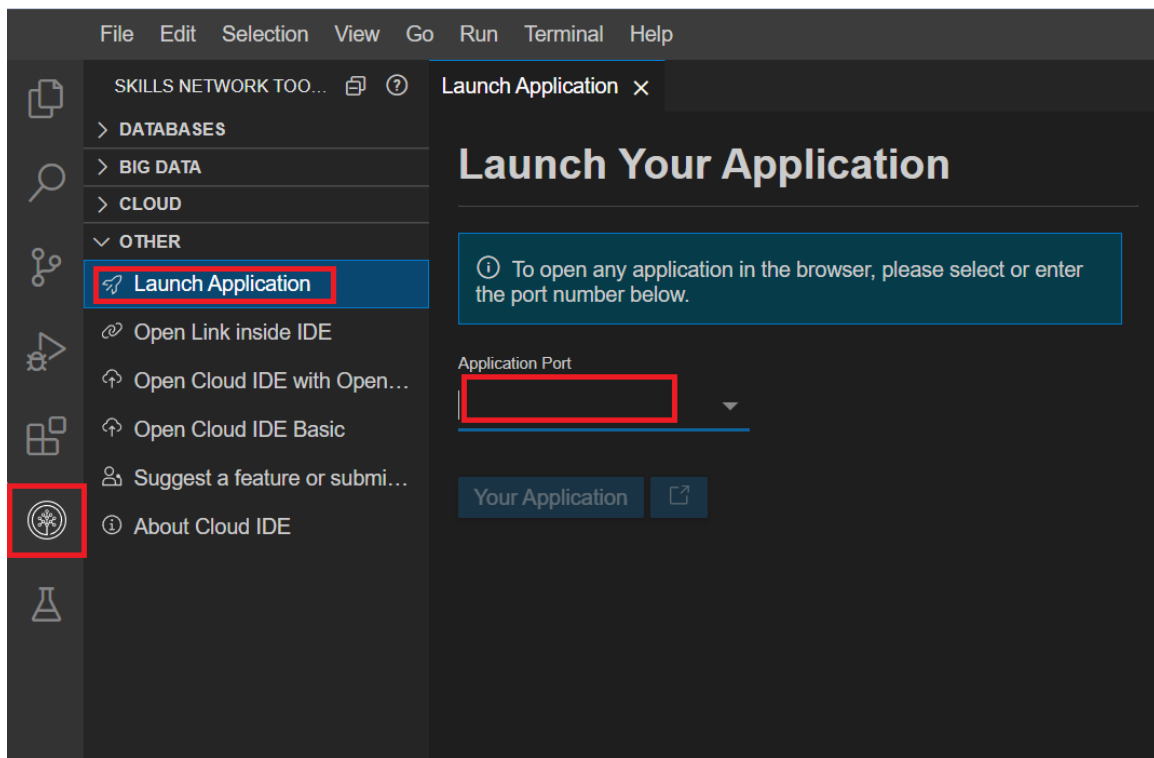
Copied!

Experiment yourself

Please have a look at the UI of the Apache Spark master and worker.

Using the *Launch Application* button you can gain access to it. Start with port 8080 first (Spark master)

Click on the Skills Network button on the left, it will open the “Skills Network Toolbox”. Then click the Other then Launch Application. From there you should be able to enter the port number as 8080 and launch.



This will take you to the admin UI of the Spark master (if your popup blocker doesn't prevent it):

Getting Started Registration Whitelisting Workday ibm w3 w3 Developer w3 Travel@IBM Launch Page Dev

Spark Master at spark://8626fea4cc2a:7077

URL: spark://8626fea4cc2a:7077

Alive Workers: 1

Cores in use: 16 Total, 16 Used

Memory in use: 61.9 GiB Total, 1024.0 MiB Used

Resources in use:

Applications: 1 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (1)

| Worker Id | Address | State | Cores |
|--|------------------|-------|---------|
| worker-20211002084753-172.18.0.3-33501 | 172.18.0.3:33501 | ALIVE | 16 (16) |

Running Applications (1)

| Application ID | Name | Cores | Memory per Executor | Resources Per |
|--------------------------------|---------------|-------|---------------------|---------------|
| app-20211002090124-0000 (kill) | pyspark-shell | 16 | 1024.0 MiB | |

Completed Applications (0)

| Application ID | Name | Cores | Memory per Executor | Resources Per Executor |
|----------------|------|-------|---------------------|------------------------|
|----------------|------|-------|---------------------|------------------------|

Please notice that you can see all registered workers (one in this case) and submitted jobs (also one in this case)

Note: The way how the lab environment works you can't click on links in the UI - in a real installation, this of course is possible.

Please repeat the steps above, re-submit the job and open the Spark worker UI by launching the application. Just use port 8081 instead of 8080 this time.

You should find your currently running job here as well.

Summary

In this lab you've learned how to setup an experimental Apache Spark cluster on top of Docker Compose. You are now able to submit a Spark job directly from python code. In the Kubernetes lab you'll learn how to submit Spark jobs from command line as well.

Changelog

| Date | Version | Changed by | Change Description |
|--------------|---------|----------------|---|
| October 2021 | 1.0 | Romeo Kienzler | Initial version |
| 01-09-2022 | 1.1 | K Sundararajan | Updated instructions for Launch Application as per new Theia IDE & the format for Changelog |
| 30-12-2022 | 1.2 | K Sundararajan | Updated pyspark installation command |